

Creating a Kubernetes cluster farm

What is MSB

- Hands-on immersive kubernetes training platform on real infrastructure
- Whole tech stack is built for the purpose of teaching Kubernetes
- Content developed by Nigel Poulton, the author of “The Kubernetes book”
- Started out of frustration with the current state of online learning



Nigel 1:32 PM

So....

2 files ▾




This is the tightest fitting M I've ever had.

Short MSB demo



Evolution of MSB



Infrastructure evolution

- MVP version (hacky version on top of Digital Ocean)
- Kubeception using VMs
- Kubeception using containerd

The MVP version

Goals

- User of Magic Sandbox should get their own cluster to train on
- The cluster should have an active terminal connection and a live visual representation of its state
- Create a proof-of-concept to gather proof points for Demo day

















Constraints

- Very limited time to implement both frontend and backend
- Limited knowledge how to actually implement disposable k8s clusters on demand
- Limited knowledge if this implementation is possible
- Limited idea of what actually Magic Sandbox does :)

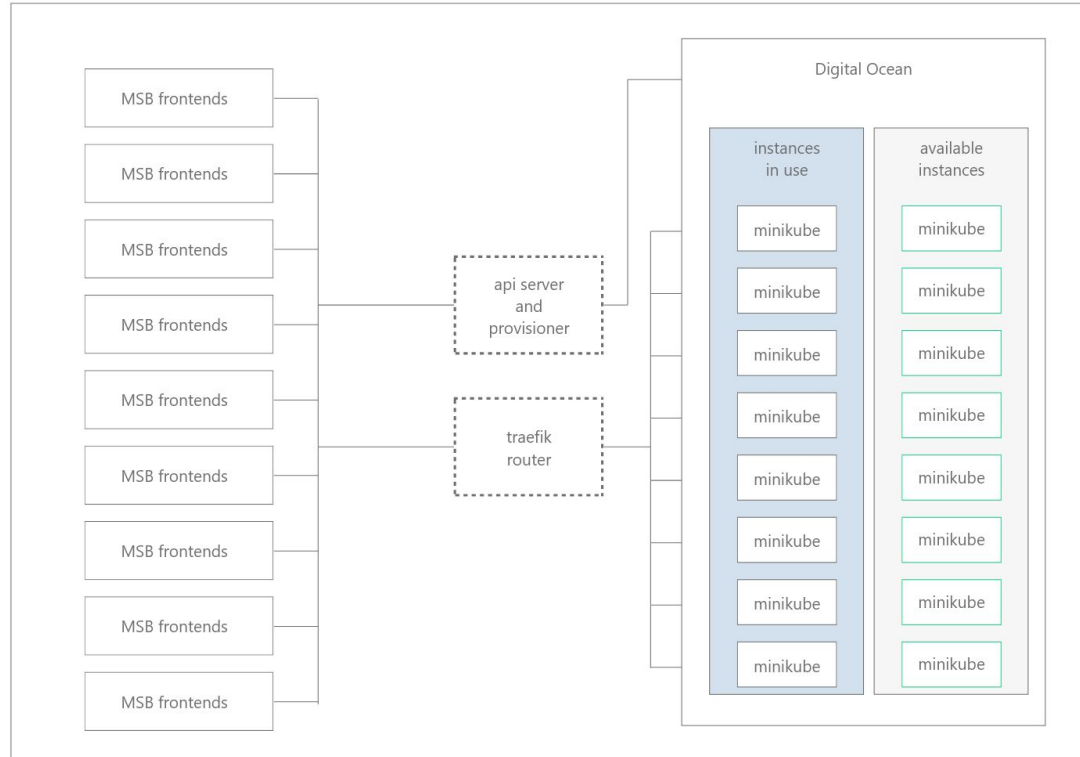
Our solution

- Generating a base Digital Ocean image with a minikube init script
- Scaling it manually
- Routing using traefik.io, then a custom stateless solution
- Using Digital Ocean labels to keep track of the clusters
- “Trivial” polling of cluster state every few seconds

Hacky Digital Ocean snapshots

 msb-beta-snap-01 Created from msb-origin	2.64 GB	 msb-origin-supervisord-nginx-ready-c... Created from msb-origin-supervisord-nginx-ready-...	4.13 GB
 msb-beta-snapshot-new Created from msb-origin	2.64 GB	 msb-origin-supervisord-cron-test-tmux Created from msb-origin-supervisord-cron-test	4.21 GB
 msb-beta-snapshot Created from msb-origin	2.64 GB	 msb-origin-supervisord-nginx-ready-h... Created from msb-origin-supervisord-nginx-ready-...	1.81 GB
 msb-origin-fix-task Created from msb-origin	1.82 GB	 msb-origin-supervisord-nginx-ready-h... Created from msb-origin-supervisord-nginx-s-2vcp...	1.8 GB
 msb-origin-rebuild-nginx Created from msb-origin	1.81 GB	 msb-origin-supervisord-nginx Created from msb-origin-before-supervisord-s-2vcp...	1.79 GB
 msb-origin-rebuild-2 Created from msb-origin	1.8 GB	 msb-origin-before-supervisord Created from msb-origin	1.74 GB
 msb-origin-rebuild Created from msb-origin	1.8 GB	 msb-origin-minimal Created from msb-origin	1.43 GB
 msb-origin-corrected Created from go-droplet	4.4 GB	 msb-origin-1527699723586 Created from msb-origin	4.38 GB

Architecture overview



Our “scaling” solution

```
// a buffered channel containing available droplets
var AvailableDroplets chan godo.Droplet = make(chan godo.Droplet, 10)

func DropletBuffer(client *godo.Client) {
    // collecting available droplets in case of a restart
    availableDroplets, _ := AvailableDropletList(context.Background(), client)

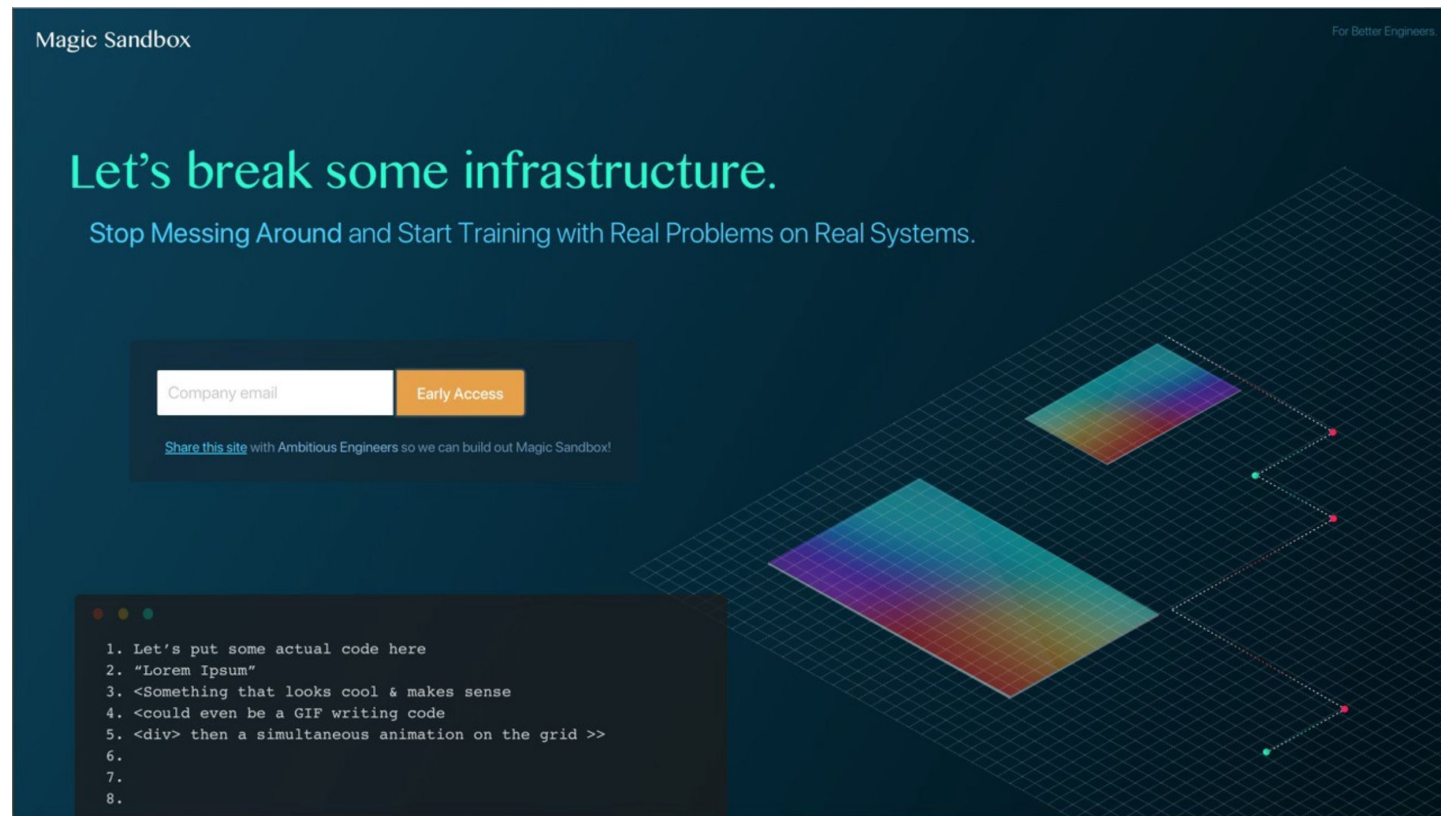
    // goroutines creating new instances and pushing into the buffered channel
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()
    go DropletGeneratorAndPusher()

    for index, droplet := range availableDroplets {
        AvailableDroplets <- droplet
    }
}
```


Problems we were encountering

- DO labeling system was failing
- Issues with traefik.io and it's bad failure scenario - ended up rewriting the router to a stateless router
- Very limited scalability - Hacker News launch exceeded our capacity in 2 minutes
- Lots of manual monitoring and interventions needed

MSB landing page at the time



MSB platform at the time



What worked well

- Least problems were with Kubernetes itself, minikube booted properly every time
- Initial MVP showed enough traction and interest to complete fundraising
- Basic ideas of MSB have been solidified
- We've realized that a lot of problems we've been having would have been handled well by Kubernetes itself

Initial Kubeception version

Goals

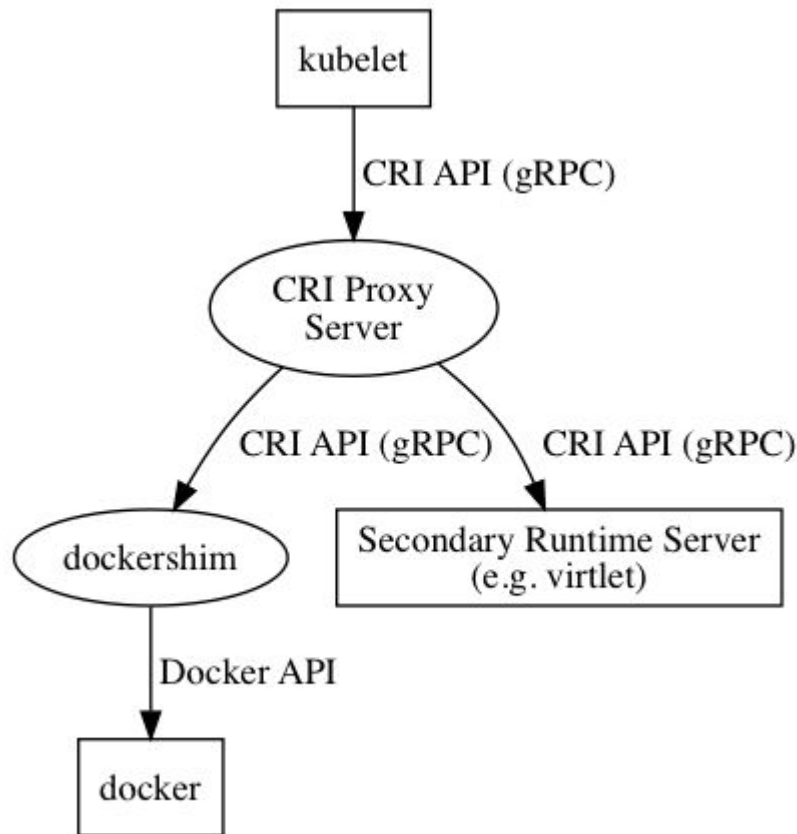
- With proper funding and team in place, build a infrastructure setup
- Infrastructure should support future MSB goals
- Create a large k8s cluster to serve smaller customer clusters
- Offload as much work to Kubernetes as possible
- Redesign frontend <-> backend communication
- More resources but bigger expectations

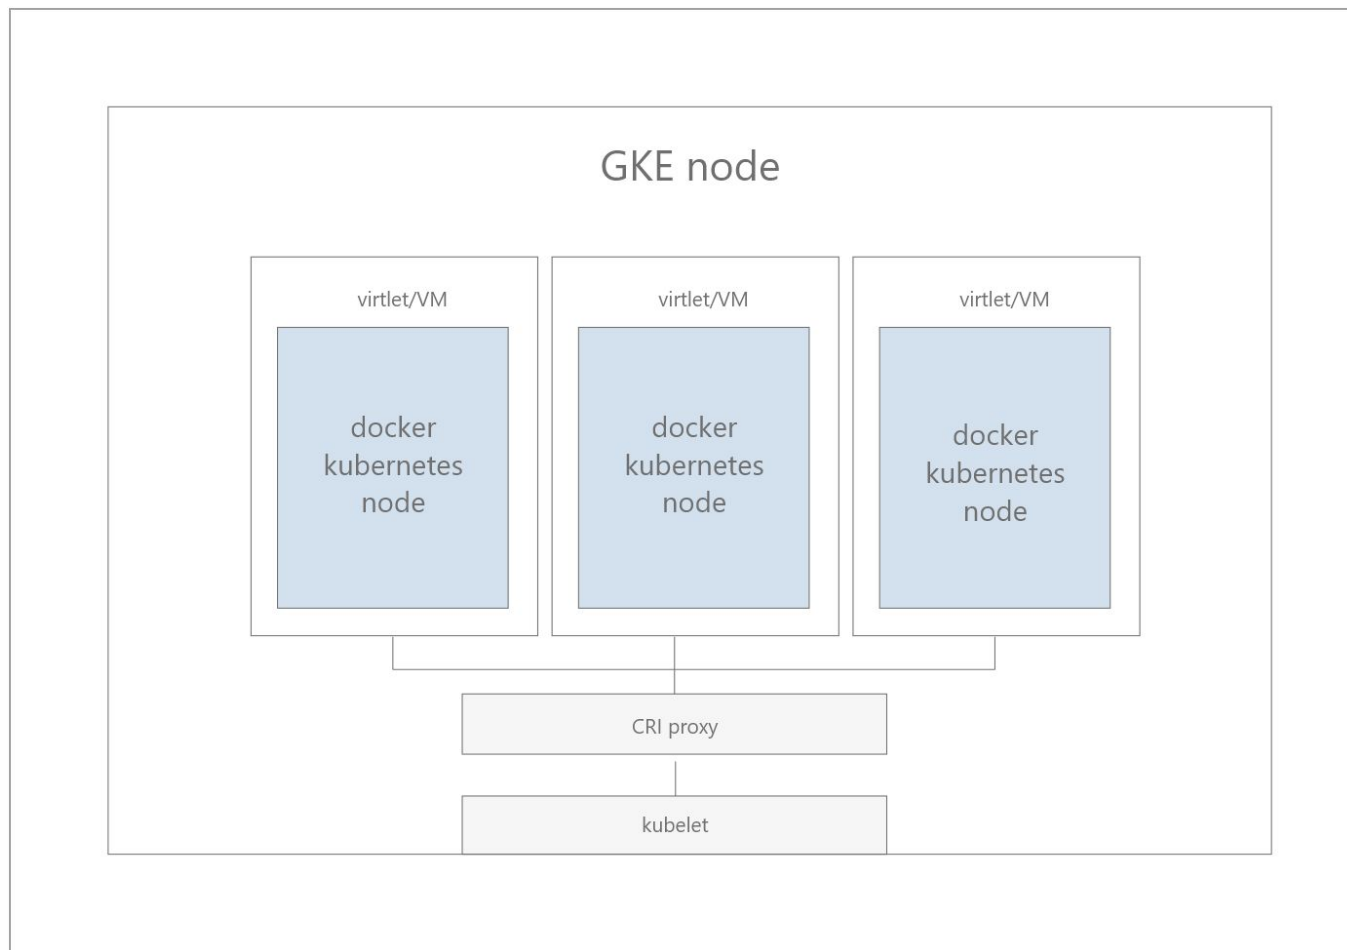
Constraints

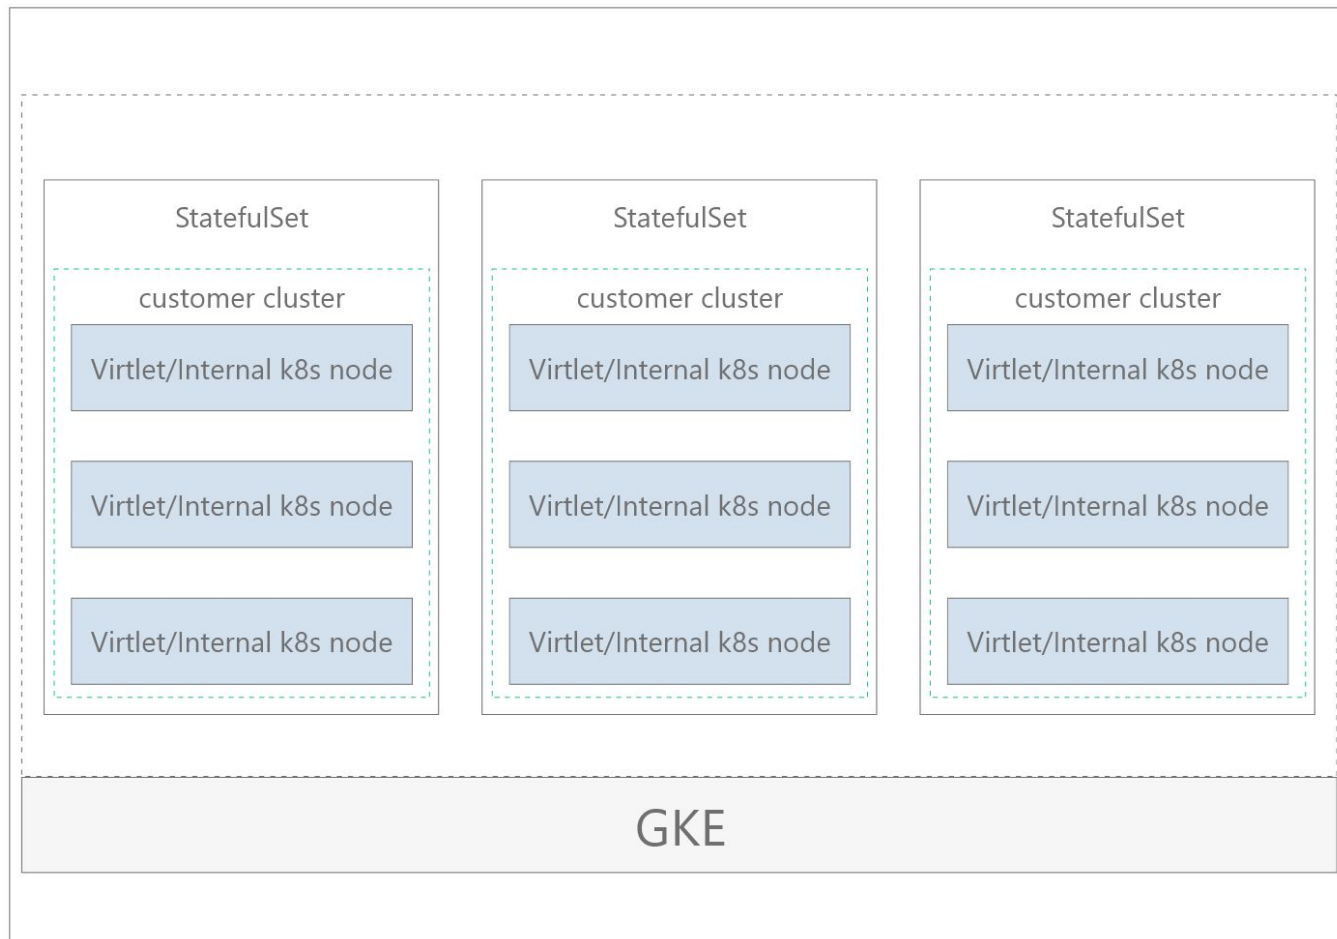
- Speed to market is again a very important driver
- Limited knowledge how to implement such a system, apart from a few blog posts
- Limited knowledge of whether large cloud providers support such “advanced” features

Solution came with virtlets

- Kubernetes CRI implementation for running VM workloads
- Container Runtime Interface (CRI) – a plugin interface which enables kubelet to use a wide variety of container runtimes, without the need to recompile
- Allowing running VMs as regular pods, with full integration with other workload objects







A custom operator to tie cluster resources together

- A custom operator ensures setup of the cluster with supporting objects (proper labels and annotations, networking, ingress)
- Ensuring there are enough “ready” clusters available in the pool
- Tracks utilization and has a percentage-based scheduling to account for spikes in usage
- Assigning a cluster to a user is a matter of changing the labels
- Build using kubernetes-sigs/kubebuilder which made implementation much easier

Frontend communication using API watchers

- Cluster state polling moved to websocket-based K8s api watchers
- Using event sourcing to recreate the state on the cluster
- Significantly reduced traffic
- All handled by traefik.io ingress controller on the main cluster without issues
- Implemented live traffic capture by hooking into docker network interfaces

Things that went well

- Again no issues came from Kubernetes itself, rather our misuse of its mechanisms
- Networking/services/ingress is rock solid
- We could finally start implementing proper monitoring and alerting
- Once all infra scripts and yamls were done, it was easy to dispose and recreate the whole “big” cluster in minutes
- GKE works as expected and gave us no issues at all
- Proper backend architecture was starting to take shape

Issues faced

- Running a docker container inside a VM inside a K8s cluster didn't feel like the best solution
- Boot up times were unnecessarily long (~3 minutes)
- VM setup not fully managed by k8s mechanisms
- Low configuration flexibility - a 3 Pod StatefulSet whose first "Pod" was hardcoded to be the master node

Switching to containerd

Goals

- Simplify architecture
- Try to get rid of VMs
- Make the clusters configurable (different number of nodes, different k8s versions...)
- Make the clusters as close to production-grade as possible (support for LoadBalancers, Ingress, PersistentVolumes, Helm support...)

Constraints

- Seamless transition for our existing users
- Not trivial to get a container runtime functional inside another container
- Try and make it the “final” rework :)

Containerd to the rescue

- A k8s node needs only 3 things:
 - container runtime
 - kubelet
 - kube-proxy
- Internal container runtime was hard to set up with Docker
- containerd allowed us to set up an internal container runtime in the container, simplifying the architecture by removing the VM

Assembling clusters on demand

- Internal operator now manages two pools of nodes - master and worker nodes, which get connected on demand
- Connection occurs when a user requests a cluster and within 10 seconds the Worker nodes are usable
- This allows per-class customization of features

Heavily dependent on Kubernetes

- Listening to kubernetes events to enable things like LoadBalancer services
- Extending the API for most introspection needs - which fits well in existing tooling
- Using community provided tooling (like the Prometheus operator) for out of the box solutions
- By choosing to rely on Kubernetes mechanisms we could easily switch cloud vendors currently

Issues faced

- Networking was a big issue, solved by switching to Cilium and BPF
 - Caused by inability to run in privileged mode
- Solving some issues required “senior-level” engineering to solve
- Not easy always to find documentation or shared experiences online for advanced issues

What we've learned

- Kubernetes is a very good and very flexible solution
- Work with k8s, not against it
- Use all available mechanisms! (init- and readinessProbes, k8s events, DaemonSets, Custom Operators...)
- Labels and annotations are very reliable and a good way to decouple systems
- **Without k8s this system would be much harder to build**, our experience was usually the bottleneck
- Frontend is also very hard
- Having a great team is the most important thing!

Thank you!

Try us out at msb.com